

# Lecture 2: Numerical Optimization for Control

(grad/SQP/QP; ALM vs. interior-point vs. penalty)

---

Arnaud Deza

August 29, 2025

ISYE 8803: Special Topics on Optimal Control and Learning

## **Overview and Big Picture of Lecture 2**

# Learning goals (what you'll be able to do)

## Goals for today

- Pick and configure an optimizer for small control problems (unconstrained & constrained).
- Derive KKT conditions and form the SQP QP subproblem for a nonlinear program.
- Explain the differences between penalty, augmented Lagrangian, and interior-point methods.

## Why?

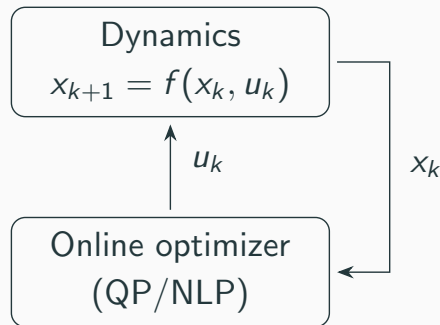
In future classes, this will help us map classic control tasks (LQR/MPC/trajectory optimization) to QPs/NLPs and choose a solver strategy.

## Roadmap for today (2 hours)

1. Big picture and some notation (5 min)
2. Unconstrained optimization: Root-finding, Newton and globalization (30 min)
3. Equality constraints: KKT, Newton vs. Gauss–Newton (20 min)
4. Inequalities & KKT: complementarity (10 min)
5. Methods: penalty  $\rightarrow$  ALM  $\rightarrow$  interior-point (PDIP) (20 min)
6. Brief look at SQP for solving hard control problems (20 min)


## Big picture: why optimization for control?

- Controller synthesis often reduces to solving a sequence of optimization problems.




## Big picture: why optimization for control?

- Controller synthesis often reduces to solving a sequence of optimization problems.
- **MPC** solves a QP/NLP online at each time step; warm-start and sparsity are critical.



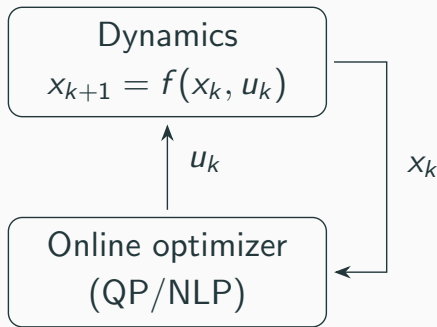
## Big picture: why optimization for control?

- Controller synthesis often reduces to solving a sequence of optimization problems.
- **MPC** solves a QP/NLP online at each time step; warm-start and sparsity are critical.
- **Trajectory optimization** (nonlinear robots) uses NLP + collocation; needs robust globalization.



## Big picture: why optimization for control?

- Controller synthesis often reduces to solving a sequence of optimization problems.
- **MPC** solves a QP/NLP online at each time step; warm-start and sparsity are critical.
- **Trajectory optimization** (nonlinear robots) uses NLP + collocation; needs robust globalization.
- **Learning-based control** backpropagates through optimizers (differentiable programming).





## Some Notation:

Given  $f(x) : \mathbb{R}^n \rightarrow \mathbb{R}$   $\frac{\partial f}{\partial x} \in \mathbb{R}^{1 \times n}$  is a row vector

This is because  $\frac{\partial f}{\partial x}$  is the linear operator mapping  $\Delta x$  into  $\Delta f$ :

$$f(x + \Delta x) \approx f(x) + \frac{\partial f}{\partial x} \Delta x$$

Similarly  $g(y) : \mathbb{R}^m \rightarrow \mathbb{R}^n$   $\frac{\partial g}{\partial y} \in \mathbb{R}^{n \times m}$  because

$$g(y + \Delta y) \approx g(y) + \frac{\partial g}{\partial y} \Delta y$$

## Some notation:

For convenience, we will define:

$$\nabla f(x) = \left( \frac{\partial f}{\partial x} \right)^T \in \mathbb{R}^n \quad (\text{column vector})$$

$$\nabla^2 f(x) = \frac{\partial}{\partial x} \left( \frac{\partial f}{\partial x} \right) = \frac{\partial^2 f}{\partial x^2} \in \mathbb{R}^{n \times n}$$

$$f(x + \Delta x) \approx f(x) + \frac{\partial f}{\partial x} \Delta x + \frac{1}{2} \Delta x^T \nabla^2 f(x) \Delta x$$

# Root-Finding

# Root Finding

Given  $f(x)$ , find  $x^*$  such that  $f(x^*) = 0$  (example: finding equilibrium of a continuous-time dynamics).

Closely related: fixed point  $f(x^*) = x^*$  (equilibrium of discrete-time dynamics)

## Fixed-Point Iteration

- Simplest solution method
- If fixed point is stable, just “iterate the dynamics” until it converges
- Only works if  $x^*$  is a stable equilibrium point and if initial guess is in the basin of attraction
- can converge slowly (depends on  $f$  i.e depending on eigenvalues)

Gradient Descent in Fixed Point Iteration applied to gradient of function  $f$ .

# Newton's Method

TLDR: Instead of solving for  $f(x) = 0$ , solve a linear for a linear approximation of  $f(x)$ .

Fit a linear approximation to  $f(x)$ :  $f(x + \Delta x) \approx f(x) + \frac{\partial f}{\partial x} \Delta x$

Set approximation to zero and solve for  $\Delta x$ :

$$f(x) + \frac{\partial f}{\partial x} \Delta x = 0 \quad \Rightarrow \quad \Delta x = - \left( \frac{\partial f}{\partial x} \right)^{-1} f(x)$$

Apply correction:  $x \leftarrow x + \Delta x$  Do this in a loop and repeat until convergence.

## Example: Backward Euler

Last time: Implicit dynamics model (nonlinear function of current state and future state)

$$f(x_{n+1}, x_n, u_n) = 0$$

Implicit Euler: this time we have  $x_{n+1}$  on the right; i.e evaluate  $f$  at future time.

$$x_{n+1} = x_n + hf(x_{n+1})$$

(Evaluate  $f$  at future time)

$$\Rightarrow f(x_{n+1}, x_n, u_n) = x_{n+1} - x_n - hf(x_{n+1}) = 0$$

Solve root finding problem for  $x_{n+1}$

- Very fast convergence with Newton (quadratic) and can get machine precision.
- Most expensive part is solving a linear system  $O(n^3)$
- Can improve complexity by taking advantage of problem structure/sparsity.

**Quick Demo of Julia Notebook: `part1_root_finding.ipynb`**

# Minimization

$$\min_x f(x), \quad f(x) : \mathbb{R}^n \rightarrow \mathbb{R}$$

If  $f$  is smooth,  $\frac{\partial f}{\partial x}(x^*) = 0$  at a local minimum.

Hence, now we have a root-finding problem  $\nabla f(x) = 0 \Rightarrow$  Apply Newton!

$$\nabla f(x + \Delta x) \approx \nabla f(x) + \frac{\partial}{\partial x}(\nabla f(x))\Delta x = 0 \quad \Rightarrow \Delta x = -(\nabla^2 f(x))^{-1} \nabla f(x)$$

$$x \leftarrow x + \Delta x$$

Repeat this step until convergence; Intuition to have about Newton:

- Fitting a quadratic approximation to  $f(x)$ ; Exactly minimize approximation



**Quick Demo of Julia Notebook: `part1_minimization.ipynb`**

## Take-away Messages on Newton

Newton is a local root-finding method. Will converge to the closest fixed point to the initial guess (min, max, saddle).

### Sufficient Conditions

- $\nabla f = 0$ : “first-order necessary condition” for a minimum. Not a sufficient condition.
- Let's look at scalar case:  $\Delta x = -\frac{1}{\nabla^2 f} \nabla f$

where: negative corresponds to “descent”,  $\nabla f$  corresponds to the gradient and  $\nabla^2 f$  acts as the “leading rate” / “step size”.

$$\nabla^2 f > 0 \quad \Rightarrow \quad \text{descent (minimization)} \qquad \nabla^2 f < 0 \quad \Rightarrow \quad \text{ascent (maximization)}$$

- In  $\mathbb{R}^n$ , if  $\nabla^2 f \succeq 0$  (positive definite)  $\Rightarrow$  descent
- If  $\nabla^2 f > 0$  everywhere  $\Rightarrow f(x)$  is strongly convex  $\rightarrow$  Can always solve with Newton

## Regularization: Ensuring Local Minimization

Practical solution to make sure we always minimize:

If  $H$  ( $H \leftarrow \nabla^2 f$ ) not positive definite, we just make it so with regularization.

While  $H \not\geq 0$ :

$$H \leftarrow H + \beta I \quad (\beta > 0 \text{ scalar hyperparameter})$$

Then do newton step as usual. I.e:

$$x \leftarrow x + \Delta x = x - H^{-1} \nabla f$$

- also called “damped Newton” (shrinks steps)
- Guarantees descent
- Regularization makes sure we minimize, but what about over-shooting?

## Line Search: Mitigating overshooting in Newton

- Often  $\Delta x$  step from Newton overshoots the minimum.
- To fix this, check  $f(x + \alpha\Delta x)$  and “back track” until we get a “good” reduction.
- Many strategies: all differ in definition of good.
- A simple + effective one is **Armijo Rule**:

Start with  $\alpha = 1$  as our step length and have tolerance  $b$  as a hyper-parameter.

$$\text{while } f(x + \alpha\Delta x) > f(x) + b\alpha\nabla f(x)^T \Delta x : \quad \alpha \leftarrow c\alpha \quad (\text{scalar } < 1 \text{ i.e } c = \frac{1}{2})$$

The intuition behind this is that  $\alpha\nabla f(x)^T \Delta x$  corresponds to the expected change in  $f$  based on a 1st order Taylor series expansion. I.e we are checking the actual decrease in  $f$  agrees with a 1st order Taylor approximation within a tolerance  $b$ .

# Constrained Optimization

## Equality-constrained minimization: geometry and conditions

**Problem.;**  $\min_{x \in \mathbb{R}^n} f(x) \quad \text{s.t.} \quad C(x) = 0, C : \mathbb{R}^n \rightarrow \mathbb{R}^m.$

**Geometric picture.** At an optimum on the manifold  $C(x) = 0$ , the negative gradient must lie in the tangent space:

$$\nabla f(x^*) \perp \mathcal{T}_{x^*} = \{p : J_C(x^*)p = 0\}.$$

Equivalently, the gradient is a linear combination of constraint normals:

$$\nabla f(x^*) + J_C(x^*)^T \lambda^* = 0, \quad C(x^*) = 0 \quad (\lambda^* \in \mathbb{R}^m).$$

**Lagrangian.;**  $L(x, \lambda) = f(x) + \lambda^T C(x).$

## A nicer visual explanation/derivation of KKT conditions

Quick little whiteboard derivation

## KKT system for equalities (first-order necessary conditions)

KKT (FOC).

$$\nabla_x L(x, \lambda) = \nabla f(x) + J_C(x)^T \lambda = 0, \quad \nabla_\lambda L(x, \lambda) = C(x) = 0.$$

**Solve by Newton on KKT:** linearize both optimality and feasibility:

$$\begin{bmatrix} \nabla^2 f(x) + \sum_{i=1}^m \lambda_i \nabla^2 C_i(x) & J_C(x)^T \\ J_C(x) & 0 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta \lambda \end{bmatrix} = - \begin{bmatrix} \nabla f(x) + J_C(x)^T \lambda \\ C(x) \end{bmatrix}.$$

*Notes.* This is a symmetric *saddle-point* system; typical solves use block elimination (Schur complement) or sparse factorizations.



Quick Demo of Julia Notebook: `part2_eq_constraints.ipynb`

# Numerical practice: Newton on KKT

## When it works best.

- Near a regular solution with  $J_C(x^*)$  full row rank and positive-definite reduced Hessian.
- With a globalization (line search on a merit function) and mild regularization for robustness.

## Common safeguards.

- *Regularize* the  $(1, 1)$  block to ensure a good search direction (e.g., add  $\beta I$ ).
- *Merit/penalty* line search to balance feasibility vs. optimality during updates.
- *Scaling* constraints to improve conditioning of the KKT system.

## Gauss–Newton vs. full Newton on KKT

**Full Newton Hessian of the Lagrangian:**  $\nabla_{xx}^2 L(x, \lambda) = \nabla^2 f(x) + \sum_{i=1}^m \lambda_i \nabla^2 C_i(x)$

**Gauss–Newton approximation:** drop the *constraint-curvature* term  $\sum_i \lambda_i \nabla^2 C_i(x)$ :

$$H_{\text{GN}}(x) \approx \nabla^2 f(x).$$

### Trade-offs (high level).

- *Full Newton*: fewer iterations near the solution, but each step is costlier and can be less robust far from it.
- *Gauss–Newton*: cheaper per step and often more stable; may need more iterations but wins in wall-clock on many problems.

**Practice tip.** Start with GN (with line search); switch to full Newton (or add low-rank updates) as feasibility improves.

# Inequality-constrained minimization and KKT

**Problem.**  $\min f(x) \quad \text{s.t.} \quad c(x) \geq 0, \quad c : \mathbb{R}^n \rightarrow \mathbb{R}^p.$

**KKT conditions (first-order).**

Stationarity:  $\nabla f(x) - J_c(x)^T \lambda = 0,$

Primal feasibility:  $c(x) \geq 0,$

Dual feasibility:  $\lambda \geq 0,$

Complementarity:  $\lambda^T c(x) = 0 \quad (\text{i.e., } \lambda_i c_i(x) = 0 \ \forall i).$

**Interpretation.**

- *Active* constraints:  $c_i(x) = 0 \Rightarrow \lambda_i \geq 0$  can be nonzero (acts like an equality).
- *Inactive* constraints:  $c_i(x) > 0 \Rightarrow \lambda_i = 0$  (no influence on optimality).

# Complementarity in plain English (and why Newton is tricky)

**What  $\lambda_i c_i(x) = 0$  means.**

- Tight constraint ( $c_i = 0$ )  $\Rightarrow$  can press back ( $\lambda_i \geq 0$ ).
- Loose constraint ( $c_i > 0$ )  $\Rightarrow$  no force ( $\lambda_i = 0$ ).

**Why naive Newton fails.**

- Complementarity = nonsmooth + inequalities ( $\lambda \geq 0$ ,  $c(x) \geq 0$ ).
- Equality-style Newton can violate nonnegativity or bounce across boundary.

**Two main strategies (preview).**

- *Active-set*: guess actives  $\Rightarrow$  solve equality-constrained subproblem, update set.
- *Barrier/PDIP/ALM*: smooth or relax complementarity, damped Newton, drive relaxation  $\rightarrow 0$ .

## **Minimization w/ Inequality Constraints**

## Three families you should know (high level)

**Goal:** Handle inequalities  $c(x) \geq 0$  (and equalities) robustly and efficiently.

### Families.

1. **Penalty:** embed violations in the objective; crank a parameter  $\rho \uparrow$ .
2. **Augmented Lagrangian (ALM):** maintain multipliers & a *moderate* penalty; solve easier subproblems.
3. **Interior-Point (PDIP):** enforce  $c(x) > 0$  via a barrier; follow the *central path* with primal–dual Newton.

**Rule of thumb.** Penalty is simplest; ALM is a strong default for medium accuracy; PDIP is the gold standard for convex QPs and very robust with Newton.

# Inequality-Constrained Minimization

**Problem Setup:**

$$\min f(x) \quad \text{s.t.} \quad c(x) \geq 0$$

**KKT conditions:**

$$\nabla f - \left( \frac{\partial c(x)}{\partial x} \right)^T \lambda = 0 \quad (\text{stationarity})$$

$$c(x) \geq 0 \quad (\text{primal feasibility}) \qquad \lambda \geq 0 \quad (\text{dual feasibility})$$

$$\lambda \circ c(x) = \lambda^T c(x) = 0 \quad (\text{complementarity})$$

**Unlike equality case, we can't directly solve KKT conditions with Newton! Why?**



## Active Set Method

- High level idea: Guess which inequalities are redundant at optimality and throw them away.
- Switch inequality constraints on/off in outer-loop and solve equality-constrained problem.
- Works well if you can guess active set well ( common in MPC where good warm-starts are common).
- Has really bad worst-time complexity.
- Usually custom heuristics are used for specific problem classes/structure.

## Lots of solution methods to use: Penalty Methods

**Penalty Method:** Replace constraints with cost terms that penalize violation!

$$\min_x f(x) + \frac{\rho}{2} \|c^-(x)\|_2^2, \quad c^-(x) := \min(0, c(x)) \text{ (elementwise).}$$

**Algorithm sketch.**

1. Start with a small  $\rho > 0$ ; minimize the penalized unconstrained objective.
2. Increase  $\rho$  (e.g.,  $\times 10$ ) and warm start from previous  $x$ .
3. Stop when  $c^-(x)$  is small enough.

**Pros.** Dead simple; reuse unconstrained machinery (Grad/Newton + line search).

**Cons.** Ill-conditioning as  $\rho \rightarrow \infty$  to satisfy constraints.; struggles to reach high accuracy; multipliers are implicit.

**Popular fix:** estimate  $\lambda$  from penalty at each iteration to be able to converge with finite  $\rho$ . Called “Augmented Lagrangian” (also closely related to ADMM).

## Augmented Lagrangian (ALM): fix penalty's weaknesses

**Core idea.** Introduce multipliers  $\lambda$  so we can keep  $\rho$  moderate and still achieve accuracy.

**Lagrangian for equality case:**  $\mathcal{L}_\rho(x, \lambda) = f(x) + \lambda^T C(x) + \frac{\rho}{2} \|C(x)\|_2^2$ .

**Outer loop.**

1.  $x^{k+1} \approx \arg \min_x \mathcal{L}_\rho(x, \lambda^k)$  (unconstrained solve).
2.  $\lambda^{k+1} = \lambda^k + \rho C(x^{k+1})$ .

**Inequalities (sketch).** Apply to the *hinge*  $c^-(x)$  and keep  $\lambda \geq 0$ :

$$\mathcal{L}_\rho(x, \lambda) = f(x) - \lambda^T c(x) + \frac{\rho}{2} \|c^-(x)\|_2^2, \quad \lambda^{k+1} = \max(0, \lambda^k - \rho c(x^{k+1})).$$

**Why it works.** Subproblems are better conditioned than pure penalty;  $\lambda$  estimates improve the model; finite  $\rho$  can reach high accuracy.

## ALM in practice (optimization loop view)

**Inner solver.** Use (damped) Newton or quasi-Newton on  $\mathcal{L}_\rho(\cdot, \lambda^k)$  with Armijo/Wolfe line search.

### **Tuning.**

- Keep  $\rho$  fixed or adapt slowly (increase if feasibility stalls).
- Scale constraints; monitor  $|C(x)|$  and stationarity.

### **When to pick ALM.**

- Nonconvex NLPs where feasibility progress matters and you want robust globalization.
- Medium accuracy is fine, or as a precursor to a polished PDIP phase on a convex QP.

Replace inequalities with barrier function in objective:

$$\min f(x), \quad x \geq 0 \quad \rightarrow \quad \min f(x) - \rho \log(x)$$

- Gold standard for convex problems.
- Fast convergence with Newton and strong theoretical properties.
- Used in IPOPT.

(Diagram:  $-\log(x)$  barrier.)

# Primal-Dual Interior Point Method

$$\min f(x) \quad \text{s.t. } x \geq 0$$

$$\rightarrow \min f(x) - \rho \log(x)$$

$$\frac{\partial f}{\partial x} - \frac{\rho}{x} = 0$$

- This “primal” FON condition blows up as  $x \rightarrow 0$ .
- We can fix this with the “primal-dual trick.”

# The Primal-Dual Trick for IPM

Introduce new variable  $\lambda = \frac{\rho}{x} \Rightarrow x\lambda = \rho$ .

$$\begin{cases} \nabla f - \lambda = 0 \\ x\lambda = \rho \end{cases}$$

- This can actually be viewed as a relaxed complementarity slackness from KKT!
- Converges to exact KKT solution as  $\rho \rightarrow 0$ .
- We lower  $\rho$  gradually as solver converges (from  $\rho \sim 1$  to  $\rho \sim 10^{-6}$ ).
- Note: we still need to enforce  $x \geq 0$  and  $\lambda \geq 0$  (with line search).

**We will use another approach from 2022 from a researcher at TRI that developed an even cooler trick.**

# Log-Domain Interior-Point Method

More general constraint case:  $\min f(x) \quad \text{s.t.} \quad c(x) \geq 0$

Simplify by introducing a “slack variable”:

$$\min_{x,s} f(x) \quad \text{s.t.} \quad c(x) - s = 0, \quad s \geq 0$$

$$\rightarrow \min_{x,s} f(x) - \rho \log(s) \quad \text{s.t.} \quad c(x) - s = 0$$

Write out Lagrangian:  $L(x, s, \lambda) = f(x) - \rho \log(s) - \lambda^T (c(x) - s)$



Apply F.O.N.C to Lagrangian from last slide:

$$\nabla_x L = \nabla f - \left( \frac{\partial c}{\partial x} \right)^T \lambda = 0$$

$$\nabla_s L = \frac{\rho}{s} + \lambda = 0 \quad \Rightarrow \quad s\lambda = \rho$$

$$\nabla_\lambda L = s - c(x) = 0$$

This second equation has a really nice interpretation: relaxed complementarity slackness

## Log-Domain Interior-Point Method

To ensure  $s \geq 0$  and  $\lambda \geq 0$ , introduce change of variables:

$$s = \sqrt{\rho}e^{\sigma}, \quad \lambda = \sqrt{\rho}e^{-\sigma}$$

Now (relaxed) complementarity is always satisfied by construction.

Plug back into F.O.N.C

$$\nabla f - \left( \frac{\partial c}{\partial x} \right)^T \lambda = 0 \qquad c(x) - \sqrt{\rho}e^{\sigma} = 0$$

We can solve these with (Gauss) Newton:

$$\begin{bmatrix} H & \sqrt{\rho}c^T e^{-\sigma} \\ c & -\sqrt{\rho}e^{\sigma} \end{bmatrix} \begin{bmatrix} \delta x \\ \delta \sigma \end{bmatrix} = \begin{bmatrix} -\nabla f + c^T \lambda \\ -c(x) + \sqrt{\rho}e^{\sigma} \end{bmatrix}$$

## Example: Quadratic Program

Super common problem to be solved in control applications: quadratic programs

$$\min_x \frac{1}{2}x^T Qx + q^T x, \quad Q \succeq 0$$

s.t.

$$Ax = b, \quad Cx \leq d$$

- Super useful in control (SQP)
- Can be solved very fast ( $\sim kHz$ ).

**Quick Demo of Julia Notebook: `part3_ipm.ipynb`**

## Penalty vs. ALM vs. PDIP: what changes?

- **Feasibility handling:**

- Penalty: encourages  $c(x) \geq 0$  via cost; feasibility only in the limit  $\rho \uparrow$ .
- ALM: balances optimality and feasibility via  $\lambda$  updates at finite  $\rho$ .
- PDIP: enforces strict interior  $c(x) > 0$ ; drives  $s_i \lambda_i = \rho \rightarrow 0$ .

- **Conditioning:**

- Penalty gets ill-conditioned as  $\rho$  grows.
- ALM keeps conditioning reasonable.
- PDIP maintains well-scaled Newton systems near the path (with proper scaling).

- **Accuracy:** Penalty (low–med), ALM (high with finite  $\rho$ ), PDIP (high; excellent for convex).

- **Per-iteration work:** Penalty/ALM solve unconstrained-like subproblems; PDIP solves structured KKT systems with slacks/duals.

# **Sequential Quadratic Programming (SQP)**

# What is SQP?

**Idea:** Solve a nonlinear, constrained problem by repeatedly solving a *quadratic program* (QP) built from local models.

- Linearize constraints; quadratic model of the Lagrangian/objective.
- Each iteration: solve a QP to get a step  $d$ , update  $x \leftarrow x + \alpha d$ .
- Strength: strong local convergence (often superlinear) with good Hessian info.

## Target Problem (NLP)

$$\min_{x \in \mathbb{R}^n} f(x) \quad \text{s.t.} \quad g(x) = 0, \quad h(x) \leq 0$$

- $f : \mathbb{R}^n \rightarrow \mathbb{R}$ ,  $g : \mathbb{R}^n \rightarrow \mathbb{R}^m$  (equalities),  $h : \mathbb{R}^n \rightarrow \mathbb{R}^p$  (inequalities).
- KKT recap (at candidate optimum  $x^*$ ):

$$\exists \lambda \in \mathbb{R}^m, \mu \in \mathbb{R}_{\geq 0}^p : \nabla f(x^*) + \nabla g(x^*)^T \lambda + \nabla h(x^*)^T \mu = 0,$$

$$g(x^*) = 0, \quad h(x^*) \leq 0, \quad \mu \geq 0, \quad \mu \odot h(x^*) = 0.$$



## From NLP to a QP (Local Model)

At iterate  $x_k$  with multipliers  $(\lambda_k, \mu_k)$ :

**Quadratic model of the Lagrangian**

$$m_k(d) = \langle \nabla f(x_k), d \rangle + \frac{1}{2} d^T B_k d$$

with  $B_k \approx \nabla_{xx}^2 \mathcal{L}(x_k, \lambda_k, \mu_k)$ .

**Linearized constraints**

$$g(x_k) + \nabla g(x_k) d = 0, \quad h(x_k) + \nabla h(x_k) d \leq 0.$$

## The SQP Subproblem (QP)

$$\begin{aligned} \min_{d \in \mathbb{R}^n} \quad & \nabla f(x_k)^T d + \frac{1}{2} d^T B_k d \\ \text{s.t.} \quad & \nabla g(x_k)^T d + g(x_k) = 0, \\ & \nabla h(x_k)^T d + h(x_k) \leq 0. \end{aligned}$$

- Solve QP  $\Rightarrow$  step  $d_k$  and updated multipliers  $(\lambda_{k+1}, \mu_{k+1})$ .
- Update  $x_{k+1} = x_k + \alpha_k d_k$  (line search or trust-region).

## Algorithm Sketch (SQP)

1. Start with  $x_0$ , multipliers  $(\lambda_0, \mu_0)$ , and  $B_0 \succ 0$ .
2. Build QP at  $x_k$  with  $B_k$ , linearized constraints.
3. Solve QP  $\Rightarrow$  get  $d_k, (\lambda_{k+1}, \mu_{k+1})$ .
4. Globalize: line search on merit or use filter/TR to choose  $\alpha_k$ .
5. Update  $x_{k+1} = x_k + \alpha_k d_k$ , update  $B_{k+1}$  (e.g., BFGS).

## Toy Example (Local Models)

**Problem:**

$$\min_{x \in \mathbb{R}^2} \frac{1}{2} \|x\|^2 \quad \text{s.t.} \quad g(x) = x_1^2 + x_2 - 1 = 0, \quad h(x) = x_2 - 0.2 \leq 0.$$

At  $x_k$ , build QP with

$$\nabla f(x_k) = x_k, \quad B_k = I, \quad \nabla g(x_k) = \begin{bmatrix} 2x_{k,1} & 1 \end{bmatrix}, \quad \nabla h(x_k) = \begin{bmatrix} 0 & 1 \end{bmatrix}.$$

Solve for  $d_k$ , then  $x_{k+1} = x_k + \alpha_k d_k$ .

## Globalization: Making SQP Robust

SQP is an important method, and there are many issues to be considered to obtain an **efficient** and **reliable** implementation:

- Efficient solution of the linear systems at each Newton Iteration (Matrix block structure can be exploited).
- Quasi-Newton approximations to the Hessian.
- Trust region, line search, etc. to improve robustness (i.e TR: restrict  $\|d\|$  to maintain model validity.)
- Treatment of constraints (equality and inequality) during the iterative process.
- Selection of good starting guess for  $\lambda$ .

# Final Takeaways on SQP

## When SQP vs. Interior-Point?

- **SQP**: strong local convergence; warm-start friendly; natural for NMPC.
- **IPM**: very robust for large, strictly feasible problems; good for dense inequality sets.
- In practice: both are valuable—choose to match problem structure and runtime needs.

## Takeaways of SQP

- SQP = Newton-like method using a sequence of structured QPs.
- Globalization (merit/filter/TR) makes it reliable from poor starts.
- Excellent fit for control (NMPC/trajectory optimization) due to sparsity and warm starts.